

# A performance evaluation using an MPI library over the communication infrastructure of Nanvix

João Fellipe Uller

`joao.f.uller@grad.ufsc.br`

Universidade Federal de Santa Catarina (UFSC) – Brazil  
Departamento de Informática e Estatística (INE)

INE410129-41000025DO/ME (20201) - Computação Paralela  
Prof. Dr. Márcio Bastos Castro

November 10, 2020



## Introduction

- Context
- Lightweight Manycores
- MPI
- Goal

## LWMPI

- Overview
- Implementation

## Evaluation

- Methodology
- Experimental Platform
- Results

## Conclusions

# Introduction

## Introduction

### Context

Lightweight Manycores

MPI

Goal

## LWMPI

Overview

Implementation

## Evaluation

Methodology

Experimental Platform

Results

## Conclusions

- Computational power vs energy consumption
- Lightweight manycores emerged to address high performance and energy efficiency demands ([FRANCESQUINI et al., 2015](#))

# Characteristics

## Introduction

Context

Lightweight Manycores

MPI

Goal

## LWMPI

Overview

Implementation

## Evaluation

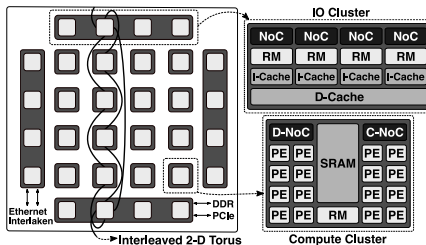
Methodology

Experimental Platform

Results

## Conclusions

- Hundreds or thousands of low-power cores on a single chip
- Heterogeneous environment
- Distributed memory system with small local memories
- Communication through message passing on a rich Network-on-Chip (NoC)



MPPA-256 Architecture

# Software Development Challenges

## Introduction

Context

Lightweight Manycores

MPI

Goal

## LWMPI

Overview

Implementation

## Evaluation

Methodology

Experimental Platform

Results

## Conclusions

- Data fetching
- Data tiling
- Asynchronous communications ([HASCOËT et al., 2017](#))
- Hand-operated routing

## Introduction

Context

Lightweight Manycores

MPI

Goal

## LWMPI

Overview

Implementation

## Evaluation

Methodology

Experimental Platform

Results

## Conclusions

### Approaches to address programmability in lightweight manycores:

#### ■ Operating Systems

- **Pros:** Bridges hardware intricacies and programmability gaps, providing **portability**
- **Cons:** Provided interface may be complex, retarding software development

#### ■ Baremetal Runtime Systems

- **Pros:** Expose rich and performance-oriented interfaces narrowed to the underlying architecture
- **Cons:** Mostly vendor-specific, resulting in non-portable software

**Duality: fast development process OR better software portability?**

# Message Passing Interface

## Introduction

Context

Lightweight Manycores

MPI

Goal

## LWMPI

Overview

Implementation

## Evaluation

Methodology

Experimental Platform

Results

## Conclusions

- A **portable** message passing standard
- Maintained and defined by the **MPI-Forum**<sup>1</sup>
- Widespread between **industry** and **academia**
- **De facto standard for message passing in HPC!**

---

<sup>1</sup>MPI-Forum website: <https://www.mpi-forum.org>

# Goal

## Introduction

- Context
- Lightweight Manycores
- MPI
- Goal**

## LWMPI

- Overview
- Implementation

## Evaluation

- Methodology
- Experimental Platform
- Results

## Conclusions

Provide a **light MPI-compliant library**, designed to cope with **architectural intricacies** of **lightweight manycores**, that is **portable** across multiple architectures and **easily extensible**.



## Introduction

- Context
- Lightweight Manycores
- MPI
- Goal

## LWMPI

- Overview
- Implementation

## Evaluation

- Methodology
- Experimental Platform
- Results

## Conclusions

LWMPI

## Introduction

Context

Lightweight Manycores

MPI

Goal

## LWMPI

Overview

Implementation

## Evaluation

Methodology

Experimental Platform

Results

## Conclusions

### LWMPI: Lightweight Message Passing Interface

- Compatible with 3.1 MPI specification (2015)
- Designed **from scratch** to be **light**
- Copes with **architectural intricacies** of lightweight manycores
- Implemented on top of a **POSIX-compliant distributed OS** (Nanvix<sup>2</sup>) to enable **portability** across different lightweight manycore architectures (PENNA et al., 2019)

---

<sup>2</sup><http://www.github.com/nanvix>

## Introduction

Context

Lightweight Manycores

MPI

Goal

## LWMPI

Overview

Implementation

## Evaluation

Methodology

Experimental Platform

Results

## Conclusions

**LWMPI currently supports the following MPI features:**

- *Runtime Management* (MPI\_Init / MPI\_Finalize)
- *Communicators*
- *Communication groups*
- *Error handlers*
- Standard *datatypes* for the C language
- Point-to-point communication (MPI\_Send and MPI\_Recv) in **synchronous mode**

# Architecture

## Introduction

Context

Lightweight Manycores

MPI

Goal

## LWMPI

Overview

Implementation

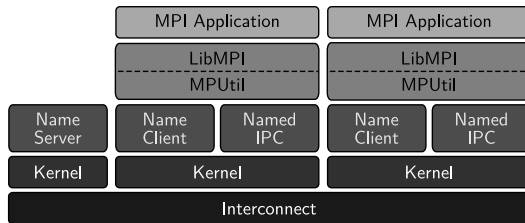
## Evaluation

Methodology

Experimental Platform

Results

## Conclusions



LWMPI Architecture

- **LibMPI**: top-level library that implements the MPI functions in a OS-independent way
- **MPUtil**: interface between LibMPI and the OS-level IPC system
  - Includes OS-dependent code

# Nanvix Support to LWMPI

## Introduction

Context

Lightweight Manycores

MPI

Goal

## LWMPI

Overview

Implementation

## Evaluation

Methodology

Experimental Platform

Results

## Conclusions

### ■ IPC Abstractions (SOUTO et al., 2020)

- Mailbox
- Portal
- Sync

### ■ Runtime Services

- Name Server
- Name Client
- Named IPC

# Communication protocol for synchronous communications

## Introduction

Context  
Lightweight Manycores  
MPI  
Goal

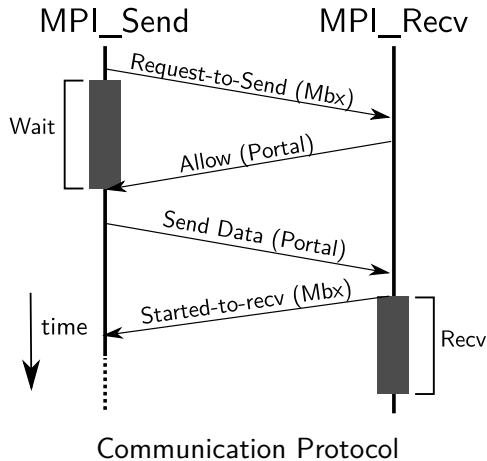
## LWMPI

Overview  
Implementation

## Evaluation

Methodology  
Experimental Platform  
Results

## Conclusions



## Introduction

- Context
- Lightweight Manycores
- MPI
- Goal

## LWMPI

- Overview
- Implementation

## Evaluation

- Methodology
- Experimental Platform
- Results

## Conclusions

# Evaluation

## Introduction

Context

Lightweight Manycores

MPI

Goal

## LWMPI

Overview

Implementation

## Evaluation

Methodology

Experimental Platform

Results

## Conclusions

### ■ Applications: CAP Bench Suite<sup>3</sup> (SOUZA et al., 2017)

- Exercise different *parallel patterns*, *task types*, *comm. intensities* and *task loads*

App	Boundary	Parallel Pattern	Comm. Intensity
FN	CPU-bound	MapReduce	Light
GF	CPU/IO Bound	Stencil	Average
KM	IO-bound	Map	Heavy

### ■ Performance evaluation

- **Target architecture:** Kalray MPPA-256 (DINECHIN et al., 2013)
- **Baseline:** vendor-specific runtime (Kalray Runtime)
- **Performance metric:** speedup

<sup>3</sup>Publicly available at: <https://github.com/nanvix/benchmarks>



# Kalray MPPA-256

## Introduction

Context  
Lightweight Manycores  
MPI  
Goal

## LWMPI

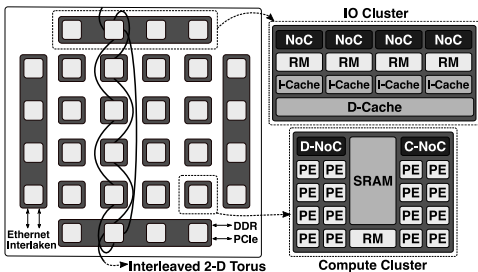
Overview  
Implementation

## Evaluation

Methodology  
Experimental Platform  
Results

## Conclusions

- 288 cores (256 GP cores + 32 firmware cores)
- 16 Compute Clusters (CCs)
- 4 I/O Clusters (IOs)
- 2 Network-on-Chips (C-NoC + D-NoC)



# Experimental Scenarios

## Introduction

Context

Lightweight Manycores

MPI

Goal

## LWMPI

Overview

Implementation

## Evaluation

Methodology

Experimental Platform

Results

## Conclusions

- CAP Bench apps have a single leader that coordinates the execution, with worker processes varying from 1 to 15 (max.)
- **Problems sizes:**
  - **FN:** numbers ranging from 1000001 to 1000129
  - **GF:** 512 × 512 image and 7 × 7 mask
  - **KM:** 30720 points and 64 centroids
- **Experimental design**
  - 30 trials for each configuration
  - Maximum CV < 1%

## Introduction

Context

Lightweight Manycores

MPI

Goal

## LWMPI

Overview

Implementation

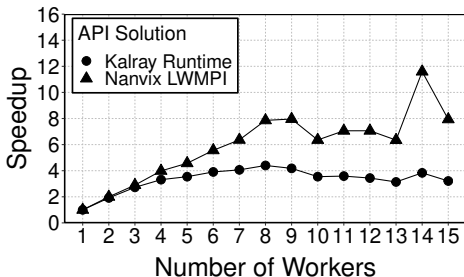
## Evaluation

Methodology

Experimental Platform

Results

## Conclusions



- Communication has little interference, since the kernel is CPU-bound
- Load imbalance with more than 8 workers
- LWMPI shows better scalability
- Easy adaptation of the kernel without significant overheads

## Introduction

Context  
Lightweight Manycores  
MPI  
Goal

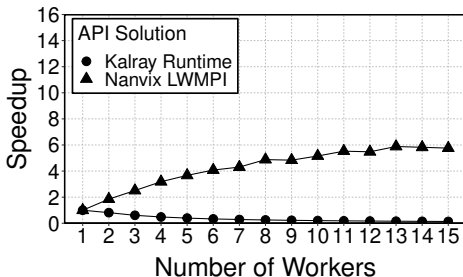
## LWMPI

Overview  
Implementation

## Evaluation

Methodology  
Experimental Platform  
Results

## Conclusions



- Small problem sizes resulted in insufficient workloads for the Kalray runtime
- Seems to be attenuated as the parallelism increases in LWMPI, proving better scalability also in these situations
- **Possibility of improvement:** asynchronous communications

## Introduction

Context  
Lightweight Manycores  
MPI  
Goal

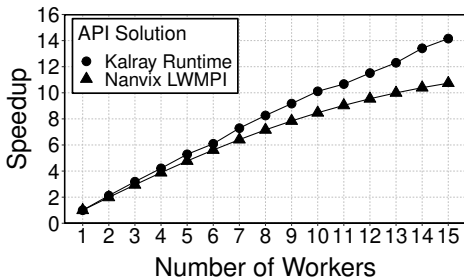
## LWMPI

Overview  
Implementation

## Evaluation

Methodology  
Experimental Platform  
Results

## Conclusions



- LWMPI and Kalray Runtime achieved similar speedups
- LWMPI showed slightly worse scalability due to coarse grain data transfers
- **Possibility of improvement:** a mechanism that dynamically chooses which OS-level comm. abstraction fits better the data transfer granularity

## Introduction

- Context
- Lightweight Manycores
- MPI
- Goal

## LWMPI

- Overview
- Implementation

## Evaluation

- Methodology
- Experimental Platform
- Results

## Conclusions

# Conclusions

## Introduction

Context

Lightweight Manycores

MPI

Goal

## LWMPI

Overview

Implementation

## Evaluation

Methodology

Experimental Platform

Results

## Conclusions

- **Our solution provides better programmability for lightweight manycores, because:**
  - It is based on an industry-standard interface (MPI)
  - It leverages a POSIX-compliant OS
  - It is portable across different lightweight manycore architectures
- **Experimental results**
  - LWMPI presents similar scalability for parallel and distributed problems, when compared to the vendor-specific runtime library (Kalray Runtime)

# References I

## Introduction

Context

Lightweight Manycores

MPI

Goal

## LWMPI

Overview

Implementation


## Evaluation


Methodology


Experimental Platform


Results


## Conclusions

 DINECHIN, B. D. de et al. A Clustered Manycore Processor Architecture for Embedded and Accelerated Applications. In: *IEEE High Performance Extreme Computing Conf.* Waltham, USA: IEEE, 2013. p. 1–6. ISBN 978-1-4799-1365-7.

 FRANCESQUINI, E. et al. On the Energy Efficiency and Performance of Irregular Application Executions on Multicore, NUMA and Manycore Platforms. *Journal of Parallel and Distributed Computing (JPDC)*, Elsevier - Academic Press, Orlando, v. 76, n. C, 2015. ISSN 0743-7315.

 HASCOËT, J. et al. Asynchronous One-Sided Communications and Synchronizations for a Clustered Manycore Processor. In: *Symp. on Embedded Systems for Real-Time Multimedia*. Seoul: ACM Press, 2017. p. 51–60. ISBN 9781450351171.

 PENNA, P. et al. On the Performance and Isolation of Asymmetric Microkernel Design for Lightweight Manycores. In: *Brazilian Symp. on Computing Systems Engineering*. Natal, Brazil: [s.n.], 2019. p. 1–8. ISSN 2324-7894.

 SOUTO, J. V. et al. Mecanismos de comunicação entre clusters para lightweight manycores no nanvix os. In: *Escola Regional de Alto Desempenho da Região Sul*. Porto Alegre, RS, Brasil: SBC, 2020. (ERAD/RS '20), p. 1–4. ISSN 2595-4164.



# References II

## Introduction

- Context
- Lightweight Manycores
- MPI
- Goal

## LWMPI

- Overview
- Implementation


## Evaluation

- Methodology
- Experimental Platform
- Results

## Conclusions



SOUZA, M. et al. Cap bench: A benchmark suite for performance and energy evaluation of low-power many-core processors. *Concurrency and Computation: Practice and Experience (CCPE)*, Wiley Online Library, v. 29, n. 4, p. 1–18, february 2017. ISSN 1532-0626.



Thank you!  
Questions?

João Felliipe Uller

`joao.f.uller@grad.ufsc.br`

Universidade Federal de Santa Catarina (UFSC) – Brazil  
Departamento de Informática e Estatística (INE)

INE410129-41000025DO/ME (20201) - Computação Paralela  
Prof. Dr. Márcio Bastos Castro

November 10, 2020



## Introduction

Context

Lightweight Manycores

MPI

Goal

## LWMPI

Overview

Implementation

## Evaluation

Methodology

Experimental Platform

Results

## Conclusions

- An Open-Source, POSIX compliant, distributed OS that targets lightweight manycores (<https://github.com/nanvix/>)
- Designed in a distributed fashion ***multikernel***
  - Multiple instances of an assymetric *microkernel*

